

Week 4 - Regular Expressions; grep; sed; awk

Regular Expressions

- many Unix utilities use regular expressions: grep, sed, awk, vi, perl, Tcl
- shell filename matches are not regular expressions (eg. *.c)
- examples in this section will use the grep utility and the file [cars](#)
- regular expressions are used to search for or match text:
- literal text can be used to search for that text

```
grep "chevy" cars
```

 - notice that "Chevy nova ..." didn't get displayed, because of the uppercase 'C'
- . matches any character (similar to ? wildcard)

```
grep ".c" cars
```

 - notice that "chevy ..." didn't get displayed, there is no character before the 'c' to match the '.'

```
grep "5..." cars
```
- [] called a character class, matches any character within the square brackets (similar to [] wildcard)

```
grep "[cC]hevy" cars
```

```
grep "[0-9][0-9][0-9][0-9][0-9]" cars
```
- [^] matches any character not within the square brackets (similar to [!] wildcard)

```
grep "[^f]ord" cars
```

 - ^ only has this special meaning if it's the first character in the square brackets
- ^ matches beginning of line

```
grep "^f" cars
```

 - ^ only has this special meaning if it's the first character in the regular expression
- \$ matches end of line

```
grep "[0-9][0-9][0-9]" cars
```

 - \$ only has this special meaning if it's the last character in the regular expression
- * following any character denotes zero or more occurrences of that character

```
grep "ford.*83" cars
```

 - . means any one single character, so .* means any number of any characters

```
grep "^^[^ ]* *^[^ ]* *65" cars
```

 - this will match '65' only in the third field
- \ inhibits meaning of special characters

```
grep "[0-9][0-9][0-9]" cars
```

 - of course, there are no records ending with a \$ sign
- regular expressions may or may not need delimiters - varies from program to program
 - eg. grep and egrep don't use delimiters, sed and awk use delimiters, usually / (forward slash)

Extended Regular Expressions

- extended regular expressions are not recognized directly by grep, can use egrep or grep -E:
- {num} following any character matches "num" occurrences of that character
egrep "[0-9]{5}" cars
- {min, max} following any character matches "min" to "max" occurrences of that character - "max" is optional
egrep "[0-9]{3,4}\$" cars
- + following any character denotes one or more occurrences of that character - same as {1,}
egrep "^^[^]+ +[^]+ +65" cars
- ? following any character denotes zero or one occurrence of that character - same as {0, 1}
egrep "ch?e" cars
- (reg-exp) parentheses used for grouping
egrep "^(^[^]+ +){2}65" cars
- | means OR, matches reg-exp on either side of the vertical bar
egrep "ford|chevy" cars

egrep "(ford|chevy) +[^]+ +65" cars
- extended regular expression characters may or may not need to be escaped - varies from program to program
 - { } + ? () | need to be escaped for use with "grep" and "sed"
 - { } + ? () | do not need to be escaped for use with "grep -E", "egrep", "sed -r", and "awk"
 - for example, the following statements have identical output:
 - grep "ford\\|chevy" cars
 - egrep "ford|chevy" cars
 - grep -E "ford|chevy" cars
 - sed -n "/ford\\|chevy/ p" cars
 - sed -nr "/ford|chevy/ p" cars
 - awk "/ford|chevy/" cars
- other examples of regular expressions
 - (Mr|Mrs) Smith - match either "Mr Smith" or "Mrs Smith"
 - Mrs? Smith - match either "Mr Smith" or "Mrs Smith"
 - [a-zA-Z]+ - match one or more letters
 - ^[a-zA-Z]*\$ - match lines with only letters
 - [^0-9]+ - match string not containing digits
 - [+]?([0-9]+[.]?[0-9]*|([0-9]+)([eE][+-]?[0-9]+)? - match valid "C" programming numbers

grep

- uses regular expression for pattern, eg. `grep 'reg-exp' filename`, then prints matched lines
- gives 0 exit status if pattern matched
- options:
 - `-c` - counts matched lines instead of printing them
 - `-i` - ignores case
 - `-n` - precedes each line with a line number
 - `-v` - reverses sense of test, eg. finds lines not matching pattern
- examples, using the file [cars](#)
 - `grep 'chevy' cars` - display only lines containing the string "chevy"
 - `grep -c 'chevy' cars` - display count of lines containing the string "chevy"
 - `grep -i 'chevy' cars` - display only lines containing the string "chevy", ignoring case
 - `grep -ic 'chevy' cars` - display count of lines containing the string "chevy", ignoring case
 - `grep -v 'chevy' cars` - display only lines not containing the string "chevy"
 - `grep -ivc 'chevy' cars` - display count of lines not containing the string "chevy", ignoring case
 - `grep -n 'chevy' cars` - display only lines containing the string "chevy", with line numbers

sed

- stream editor
- `sed 'address instruction' filename`
- checks for address match, one line at a time, and performs instruction if address matched
- prints lines to standard output by default (supressed by `-n` option)
- addresses
 - can use a line number, to select a specific line (for example: 5)
 - can specify a range of line numbers (for example: 5,7)
 - can specify a regular expression to select all lines that match
 - default address (if none is specified) will match every line
- instructions
 - `p` - print line(s) that match the address (usually used with `-n` option)
 - `d` - delete line(s) that match the address
 - `q` - quit processing at the first line that matches the address
 - `s` - substitute text to replace a matched regular expression, similar to vi substitution
- examples, using the file [cars](#)
 - `sed '3,6 p' cars` - display lines 3 through 6 (these lines will be doubled, since all lines printed by default)
 - `sed -n '3,6 p' cars` - display only lines 3 through 6
 - `sed '5 d' cars` - display all lines except the 5th
 - `sed '5,8 d' cars` - display all lines except the 5th through 8th
 - `sed '5 q' cars` - display first 5 lines then quit, same as `head -5 cars`
 - `sed -n '/chevy/ p' cars` - display only lines matching regular expression, same as `grep 'chevy' cars`
 - `sed '/chevy/ d' cars` - delete all matching lines, same as `grep -v 'chevy' cars`
 - `sed '/chevy/ q' cars` - display to first line matching regular expression
 - `sed 's/chevy/gm /' cars` - substitute "chevy" with "gm " on each matching line
 - `sed 's/[0-9]/*/' cars` - substitute first occurrence of a digit on each line with an asterisk
 - `sed 's/[0-9]/*g' cars` - substitute every occurrence of a digit on each line with an asterisk
 - `sed -r '5,8 s/[0-9]+/***/' cars` - substitute only on lines 5 to 8
 - `sed -nr '5,8 s/[0-9]+/***/ p' cars` - substitute only on lines 5 to 8, print only matched lines
 - `sed -r '/ford/ s/[0-9]+/***/' cars` - substitute only on lines containing "ford"

- `sed -nr '/ford/ s/[0-9]+/***/ p' cars` - substitute only on lines containing "ford", print only matched lines
- `sed -r 's/[0-9]+/*** & ***/' cars` - & is the value of the string matched by reg-exp
- `sed -r 's/^[0-9]*([0-9]+).*/The first number is \1/' cars` - \1 is the string matched within the first group
- `sed -r 's/([]+)+([]+)+/\2\1/' cars` - swap first two fields
- `sed -r 's/([]+)+([]+)+([]+)+([]+)+([]+)/We have a \1 \2 at only $5/' cars`
- `sed -nr '/ford/ s/[]+ +([]+)+ []+ + []+ +([]+)/We have an amazing \1 for the low price of $2! What a steal!/ p' cars`
- `sed -nr '/$1/' s/[]+ +([]+)+ []+ + []+ +([]+)/We have an amazing \1 for the low price of $2! What a steal!/ p' cars`
- when using multiple commands, the following statements have identical output:
 - `sed 's/ford/Ford/' cars | sed 's/chevy/Chevy/'`
 - `sed -e 's/ford/Ford/' -e 's/chevy/Chevy/' cars`
 - `sed 's/ford/Ford/; s/chevy/Chevy/' cars`

awk

- pattern matching and processing
- `awk 'pattern {action}' filename`
- checks for pattern match, one line at a time, and performs action if pattern matched
- variables
 - NR is an awk variable meaning the line number of the current record
 - NF is an awk variable meaning the number of fields in the current record
 - \$n are awk variables, meaning the value of the nth field (field delimiter is space or tab)
 - \$0 is the entire record
 - IFS is an awk variable specifying the input field separator, defaults to space or tab
 - can also be specified using the -F option
 - OFS is an awk variable specifying the output field separator, defaults to a space
 - user-defined variables don't need to be declared, any unquoted string is assumed to be a variable
 - variables are automatically initialized, 0 if used numerically, null if used as a string
- numeric comparison will be done if both sides are numeric (eg. `$3 > 65`), otherwise string comparison will be done
- pattern
 - can use a line number to select a specific line, by comparing it to NR (for example: `NR == 2`)
 - can specify a range of line numbers (for example: `NR == 2, NR == 4`)
 - can specify a regular expression, to select all lines that match
 - can compare field values to literals or variables (for example: `$3 == 65`)
 - can check for a regular expression match within a field by using the ~ operator (for example: `$2 ~ /[0-9]/`)
 - BEGIN is a special pattern that causes execution of the action before any records have been read
 - usually used to initialize variables, print header lines for reports, etc.
 - END is a special pattern that causes execution of the action after all records have been read
 - usually used to calculate totals and averages, print summary lines for reports, etc.
 - every line is selected if no pattern is specified

- instructions
 - print - print line(s) that match the pattern, or print fields within matching lines
 - print is default action if no action is specified
 - printf - a C-like version of print, with similar format specifications and no automatic new-line
 - there are many, many instruction, including just about all C statements with similar syntax
- examples, using the file [cars](#)
 - `awk 'NR == 2, NR == 4' cars` - display the 2nd through 4th lines (default action is to print entire line)
 - `awk '/chevy/' cars` - display only lines matching regular expression, same as `grep 'chevy' cars`
 - `awk '{print $3}' cars` - display third field of all lines
 - `awk '{print $3 $1}' cars` - display third and first field of all lines
 - `awk '{print $3, $1}' cars` - includes an output field separator (variable OFS) because of the comma
 - `awk -F:' ' '{print $6}' /etc/passwd` - specifies that : is input field separator, default is space or tab
 - `awk '/chevy/ {print $3, $1}' cars` - display third and first field of lines matching regular expression
 - `awk '$3 == 65' cars` - display only lines with a third field value of 65
 - `awk '$5 <= 3000' cars` - display only lines with a fifth field value that is less than or equal to 3000
 - `awk '$5 <=$price' '{print $1, $2, $5}' cars` - \$price is a shell variable, not an awk variable, e.g. first execute: `price=3000`
 - `awk '$5 > 3000 && $5 < 9000' cars` - display lines where 5th field is in range 3000 to 9000
 - `awk '$5 < 3000 || $5 > 9000' cars` - display lines where 5th field is outside of range 3000 to 9000
 - `awk '$2 ~ /[0-9]/' cars` - searches for reg-exp (a digit) only in the second field
 - `awk '{printf "%-30s%20s\n", $5, $1}' cars` - display 5th field left-justified in a 30 character field, 1st field right-justified in a 20 character field
 - `awk '$5 >=$price' '$5 = $5 * 0.9' '{print}' cars` - if field 5 >= shell variable \$price then reduce field 5 by 10%, e.g. first execute: `price=$5000`
 - `awk '$3 < "8"' cars` - double quotes force string comparison
 - `awk '$1 < "honda"' cars` - double quotes force string comparison, otherwise "honda" would be treated as a variable
 - `awk 'NF != 5' cars` - display lines without 5 fields
 - `awk 'BEGIN {OFS="~"} {print $1, $2}' cars` - display 1st and 2nd field of each record, separated by ~
 - `awk '{OFS="~"; print $1, $2}' cars` - same result, but much less efficient
 - `awk '{printf "%-8s%-8s%-8s%-8s%-8s\n", $2, $1, $3, $4, $5}' cars` - swap first two fields
 - `awk '{printf "We have a %-8s %-8s at only $%s\n", $1, $2, $5}' cars`
 - `awk '/ford/ {print "We have an amazing " $2 " for the low price of $" $5 "! What a steal!"}' cars`
 - `awk '/$1/' '{print "We have an amazing " $2 " for the low price of $" $5 "! What a steal!"}' cars`